

Image compositing for photographic lighting

Tjaart van der Walt

May 6, 2015

Contents

1	Introduction	2
1.1	Credit	2
1.2	Rationale	2
2	Installation	3
2.1	Source code	3
2.2	Dependencies	3
2.2.1	Python interpreter	3
2.2.2	OpenCv	3
2.3	Python	3
2.3.1	Pip	3
2.4	Installation methods	4
3	Test data	4
3.1	Input images	4
3.1.1	Basket	4
3.1.2	Cafe	5
4	Usage	5
4.1	Command line options and arguments	5
4.2	Basis lights	6
4.2.1	Fill light	6
4.2.2	Edge light	6
4.2.3	Diffuse color light	6
4.3	Light modifiers	6
4.3.1	Per-Object modifier	7
4.3.2	Soft light modifier	7
4.3.3	Regional lighting modifier	7

5 Implementation	7
5.1 Basis lights	7
5.1.1 Fill light	7
5.1.2 Edge light	9
5.1.3 Diffuse color light	9
5.2 Light modifiers	11
5.2.1 Per-Object modifier	11
5.2.2 Soft light modifier	13
5.2.3 Regional lighting modifier	14
6 Future work	16

1 Introduction

1.1 Credit

This project is an implementation of the concepts found in the paper [User-assisted Image Compositing for Photographic Lighting¹](#) by Ivaylo Boyadzhiev, Sylvain Paris and Kavita Bala. All the mathematical models found in this document is taken from the paper. The contribution of this project is that it provides a Free Software implementation of these concepts.

1.2 Rationale

We want to provide better software support for a new workflow that has been emerging in the area of photography of static scenes. This workflow differs from the traditional workflow in the fact that it does not require the photographer to meticulously consider his lighting setup before taking a photo. With the new workflow the photographer takes multiple photos of the scene from a fixed vantage point² with a moving light source. Photographers currently using this workflow then exports these images into an image manipulation software package and manually edits the input images into one well lit output image.

This project provides support to automate a lot of the work that can currently only be done by image manipulation experts. We provide 6 command line scripts³ to attain this goal.

Light type	Description
Fill light	Provides even illumination across the image
Edge light	Emphasizes the main edges in the scene
Diffuse color light	Emphasizes the base colors of objects
Per object modifier	Emphasizes a particular object
Soft light modifier	Removes harsh shadows and highlights
Regional lighting modifier	Balances the overall lighting of the scene

¹http://www.cs.cornell.edu/projects/light_compositing/

²Typically a tripod.

³For detailed usage see the [Usage](#) section.

2 Installation

2.1 Source code

The source for the package is available on [GitHub](https://github.com/tjaartvdwalt/light-compositing)⁴

2.2 Dependencies

2.2.1 Python interpreter

To use this project you need Python2 (preferably 2.7+) installed on your system. It can be downloaded from [here](#).⁵

On Arch Linux you can install Python using the package manager.

```
pacman -S python2
```

In addition to the Python interpreter the following Python packages will need to be installed.

- NumPy
- SciPy

There should however be no need to install them manually, as they should get built by the installation script.

2.2.2 OpenCv

You will need a working OpenCv2 (preferably 2.4+) installation, with the Python bindings installed. You can download it from [here](#).⁶

On Arch Linux you can install OpenCv using the package manager.

```
pacman -S opencv
```

2.3 Python

Our system has been developed using Python2. We have used some Python2 specific syntax in various places in our code⁷. At some point we would like to refactor the code to make it Python3 compatible.

2.3.1 Pip

`pip` is a command line tool that allows you to install packages from the [Python Package Index](#). The tool handles all dependency management, making it a very appealing way to install Python packages.

⁴<https://github.com/tjaartvdwalt/light-compositing>

⁵<https://www.python.org/downloads/>

⁶<http://opencv.org/downloads.html>

⁷Simply because we did not know any better at the point of writing.

If you have Python 2.7.9 or later installed, the `pip` command is included with the interpreter. Otherwise you can install `pip` using [this⁸](#) script.

```
python get-pip.py
```

2.4 Installation methods

There are two ways to easily install the `light-compositing` package. The first is to use the Python `Distutils` directly. For this you will have to download the source and execute the `setup.py` script in the root directory of the project using the command:

```
python setup.py install
```

The second way is to install it from the `Python Package Index`.

```
pip install light-compositing
```

Both of these installation methods installs the `light_compositing` package into the Python `site-packages`⁹ directory. On a Unix based system it will also install the executable scripts into the `/usr/bin` directory.

3 Test data

We have included our test data in the `test_data` sub directory of our submission. This is a subset of the images used by the authors of the paper that this implementation is based on. We do not own the copyright of these images, but they appear to be in public domain. The original images can be downloaded from [here](#).¹⁰

3.1 Input images

We chose two sets of input images from the available test data to do our development.

3.1.1 Basket

The basket of fruit gives us a typical use case of a static scene. The basket contains several distinct objects providing a good baseline for our edge light model. The shadows cast around the base of the bowl also provide a good test for soft lighting modifier. This set consists of 128 images. In figure 1 we see a few representative images from this set.

⁸<https://bootstrap.pypa.io/get-pip.py>

⁹Located at `/usr/lib/python2.7/site-packages/light_compositing-0.0.0-py2.7.egg` on my system.

¹⁰http://www.cs.cornell.edu/projects/light_compositing/download/light_compositing_input_data.zip



Figure 1: Representative images from the basket of fruit set. This set consists of 128 images in total

3.1.2 Cafe

In contrast to the basket image set, the cafe provides a much larger scene, so that a much smaller area of the scene is lit in any one image. This is a good test for the fill light model, because the input light will be less uniform than with the basket images. It also gives us the best opportunity to test our regional lighting modifier. The red chair in the cafe is used in the paper as the baseline for diffuse color light. This image set consists of 113 images. In figure 2 we see a few representative images from the set.



Figure 2: Representative images from the cafe set. This set consists of 113 images in total

4 Usage

Currently our application does not contain a fully fledged graphical user interface. Each light type is a script that can be run from the command line.

4.1 Command line options and arguments

Most of the scripts takes a `directory` argument. This is the directory where the input images are found. We assume that these images are named in the format `000.png ... xyz.png`. If no `--count` option is given, these image files should be the only files in the directory. ¹¹

¹¹We use the number of files in the directory as the default value for `COUNT`.

Most of the scripts have a `--count` option, that allows you to set how many of the input images should be used for the calculation.¹². This is useful if you do not want to waste time doing calculations on all the input images. From the results in the paper it seems that setting `COUNT > 15` will yield good results in most cases.

All of the scripts have a `--verbose` option. Setting this option prints debug information.

All of the scripts have a `--output` option. You can use this option to set the name of the output file you want to generate. If left unset a file name will be based on the light type that was used to generate it. For example running the `fill_light` script will produce `fill_light.png`

The edge light, and diffuse color light scripts have a `--downsample` option. This option value determines how many times the input images will get down sampled before being passed to the optimization algorithm. This reduces the time required to find an optimized result dramatically. The resulting lambdas are then applied to the original sized images. This is very useful if you want to get a pretty good solution in quick time.

4.2 Basis lights

Depending on the size and number of input images used it can take a long time to compute the basis lights. In the next sections we show the commands used to generate some of the results in this report. That should give the reader a good understanding of the options available to each script.

4.2.1 Fill light

```
fill_light --verbose --count 10 test_data/basket
```

4.2.2 Edge light

```
edge_light --verbose --downsample 3 --count 10 test_data/cafe
```

4.2.3 Diffuse color light

```
diffuse_color_light --verbose --downsample 3 --count 10 test_data/basket
```

4.3 Light modifiers

The light modifiers can be calculated in real time. With further development these scripts could be incorporated into a graphical user interface that allows a user to manipulate the lighting of his image.

¹²Unfortunately, this setting always uses the first `count` images in the directory, so you cannot rerun the program with a set of randomly selected input images at this stage.

4.3.1 Per-Object modifier

The per object modifier takes the output files generated by the `fill_light`, `edge_light` and `diffuse_color_light` scripts as arguments. If these arguments are not specified it uses the default values for each of these lighting types.

The per object modifier has a `--mask` option that gives the path to a binary image that defines the object of interest. In addition it has `--fill` `--edge` and `--diffuse` options that gives the weight of each lighting type for this modifier.

In this example we give each light type equal weight.

```
object_modifier --mask test_data/basket_regions/corn.png \
> --fill 1 --edge 1 --diffuse 1 fill_light.png \
> edge_light.png diffuse_color_light.png
```

4.3.2 Soft light modifier

The soft light modifier has a `--sigma` option that defines the σ value. This value controls how much a particular light is diffused as described in the implementation section.

```
soft_light_modifier --sigma 0.5 test_data/basket
```

4.3.3 Regional lighting modifier

The regional light modifier takes an input image as argument instead of an input directory like most of the other lighting types. In addition it has a `--beta` option that defines the β value. When this value is positive, it emphasizes the bright areas in the input image, and when it is negative it emphasizes dark areas in the input image.

```
regional_light_modifier --beta 0.5 my_image.png
```

5 Implementation

5.1 Basis lights

5.1.1 Fill light

Fill light tries to give even illumination everywhere in the input image. We assume that our input lights will be approximately uniformly distributed across the input images.

Model Our fill light is defined by:

$$I_{\text{fill}}(p) = \frac{\sum_i w_i(p) I_i(p)}{\sum_i w_i(p)}$$

Where

$$w_i(p) = \frac{\bar{I}_i(p)}{\bar{I}_i(p) + \epsilon}$$

$$\bar{I}_i(p) = I_i(p) \cdot (0.2990, 0.5870, 0.1140)$$

$$\epsilon = 0.01$$

Results In figure 3 we compare our fill light with an unweighted average of all 128 input images. We note that there is not much difference between the two images.



Figure 3: Fill light applied to the bowl of fruit image set.

Left: Average light, **Right:** Fill light

To illustrate the real value of the weight component in our fill light, we chose 30 of the cafe images in which the right side was more prominently lit than the left side. In figure 4 we can see that the unweighted average performs poorly, but the fill light gives a good result.



Figure 4: Fill light applied to the cafe image set. The images were selected in such a way that the left side is left relatively dark.

Left: Average light, **Right:** Fill light

5.1.2 Edge light

Edge light tries to emphasize the main edges in a scene. We leverage the fact that the main edges of the scene will always be located in the same place in the image, whereas occasional shadows and highlights will be sporadically distributed over the input images.

Model Our edge light is defined as:

$$I_{\text{edge}} = \sum_i \lambda_i I_i$$

$$\arg \min_{\{\lambda_i\}} \sum_p h(p) \left| \nabla \left(\sum_i \lambda_i I_i(p) \right) - G(p) \right|^2$$

Where G is the gradient map of the combined input images. $h(p)$ is the per pixel weight designed to give more weight to pixels with a well defined orientation.

Implementation notes We use a Sobel mask to calculate the gradients orientations and magnitudes for each input image. From this we calculate a histogram of orientations at each pixel across all the input images. We create bins for every 5° . From this histogram we choose the bin with the most entries. From this bin we choose the entry with the greatest magnitude, and add it as an entry to our gradient map $G(p)$ ¹³. Furthermore, we normalize our histogram, and add the value of the largest bin to the weight map $h(p)$.

We then minimize the given equation using `scipy.optimize.minimize`. We use the truncated Newton method, and we limit the values to be between $(0, 1)$. To obtain quicker results one might consider using the `--downsample` option. When this is enabled we calculate the λ values on a downsampled version of the input images, but apply the result to the full size images.

Results In figure 5 we compare the results of our edge light with that of the average light for the first 10 images in our cafe image set. We notice that the images where the edges are more pronounced are assigned greater weight. For this set of input images 009.png gets the highest weight.

5.1.3 Diffuse color light

Diffuse color light tries to emphasize the base color of objects.

¹³We add the orientation to our gradient map, not the magnitude.



Figure 5: Edge light applied to 10 of the cafe set of images.
Left: Average light, **Right:** Edge light

Model For this model we handle colorful objects and neutrally colored objects separately. For colorful objects we define a function that favors diffuse reflection instead of the specular reflection. It is a key insight that the specular component appears white, so the stronger it is the less colorful the object appears.

$$\mathbf{I} = d\mathbf{D} + s\mathbf{W}$$

The saturation angle is given by $\angle(I, W) = \arccos\left(\frac{I}{\|I\|} \cdot \frac{W}{\|W\|}\right)$ where $W = (1, 1, 1)^T$,

We notice that for a fixed value d the angle decreases as s increases. We want to choose λ 's that maximizes the angle between the pixels in our input images and the specular component.

$$\arg \max_{\{\lambda_i\}} \sum_p \hat{w}(p) \angle \left(\sum_i \lambda_i I_i(p), W \right) \quad (1)$$

Similarly to the fill light, we apply weights to discourage the selection of dark pixels.

$$\hat{w}(p) = \frac{\sum_i \lambda_i I_i(p)}{\sum_i \lambda_i I_i(p) + \epsilon}$$

For neutrally colored objects we define a function that encourages similarity between the average image and our result.

$$\arg \min_{\{\lambda_i\}} \sum_p \angle \left(\sum_i \lambda_i I_i, \frac{1}{N} \sum_i I_i \right) \quad (2)$$

To make sure we apply this only to neutral colors, we define a balancing term:

$$\alpha(p) = \exp\left(\frac{-\angle\left(\frac{1}{N}\sum_i I_i(p), W\right)^2}{2\sigma^2}\right)$$

By combining equation 1 and equation 2 we define our diffuse color light as:

$$I_{\text{diffuse}} = \sum_i \lambda_i I_i$$

where

$$\begin{aligned} \arg \min_{\{\lambda_i\}} \sum_p & \left[\alpha(p) \angle \left(\sum_i \lambda_i I_i(p), \frac{1}{N} \sum_i I_i(p) \right) \right. \\ & \left. - (1 - \alpha(p)) \hat{w}(p) \angle \left(\sum_i \lambda_i I_i(p), W \right) \right] \end{aligned}$$

Implementation notes The diffuse color light shares the edge light optimization code. There might be a few bugs lurking in the implementation, as we cannot reproduce some of the results published in the paper. For a detailed explanation see the results section below.

Results In figure 6 we compare the fill light with the diffuse color light for the first 10 cafe images. The results are less than satisfactory, as our diffuse color light looks similar to the average light. My current theory is that the number average colored pixels are overpowering bright colored pixels, leading to the average image.¹⁴

In figure 7 we compare the fill light with the diffuse color light for the first 10 basket images. In these images we can see slight improvements to the color of the banana and pepper. The table in the background is also brightened.

5.2 Light modifiers

5.2.1 Per-Object modifier

The per-object modifier controls the spread of light. It is inspired by photographic equipment like a snoot.

Model

- We compute the basis lights as discussed in the previous section.
- We apply them only to pixels within a selected object.

¹⁴In the paper they appear to apply the diffuse light only to the brightly colored chair in the cafe image so that bright pixels carry most weight. Unfortunately at this stage we have no mechanism to easily emulate these results.



Figure 6: Diffuse color light applied to the first 10 images in the cafe set.
Left: Average light, **Right:** Diffuse color light



Figure 7: Diffuse color light applied to the first 10 images in the basket set.
Left: Average light, **Right:** Diffuse color light

- The users can then mix these colors interactively using 3 track bars.
- To ensure smooth blending with the rest of the image we apply a bilateral filter

Implementation notes For our implementation, objects are defined by a mask image. This image contains white pixels in the object area, and black pixels in the background area. If we had time to create a GUI we would have allowed the user to define the object using a pointing device.

Our current implementation lacks the ability to blend the pixels of the masked image back into the resulting image. This is clearly apparent in our results.

Results In figure 8 we apply the per object modifier on the red chair in the cafe image. We create the object by giving equal weight to the fill light, edge light and diffuse light, and then we merge the result back into the fill light image.



Figure 8: Per object modifier applied to the bright red chair.
Left: Fill light, **Right:** Soft light modifier

5.2.2 Soft light modifier

The soft light modifier reduces the effect of harsh shadows and highlights. It is inspired by photographic equipment like umbrellas and soft boxes

Model Our soft light modifier is defined by:

$$I_{\text{soft}} = \sum_i \lambda_i I_i$$

where

$$\text{soft}_{\sigma_S}(\Lambda) = \frac{||\Lambda||}{||S\Lambda||} S\Lambda$$

$$\Lambda = (\lambda_1, \dots, \lambda_N)^T$$

$$S_{ij} = \exp\left(\frac{-||I_i - I_j||^2}{2\sigma_s^2}\right)$$

Implementation notes The soft light modifier does not maintain the image intensity. To fix this we do a global rescaling of the image intensity. In order to do that we convert the image into the HSV colorspace. A limitation of our implementation is that we do not know how much to rescale the image. Currently this is a hardcoded and does not work well for all images.

Results In figure 9 we compare the fill light, and soft light modifier. The biggest changes we can notice is that the reflection from the table is removed by the soft light modifier.



Figure 9: Soft light modifier applied to the bowl of fruit images.
Left: Fill light, **Right:** Soft light modifier

5.2.3 Regional lighting modifier

The regional lighting modifier balances the light intensities in a scene to emphasize (or de-emphasize) a specific region in a scene. It provides a way to adjust the light balance across the scene at a coarse level.

Model

- We are only interested in balancing the lighting, so we work with the intensity images \bar{I}_i
- We calculate the first component using Principle Component Analysis. This captures the different areas of illumination
- We translate our image to the log domain $P = \ln(\bar{I}_i)$
- We ensure the overall image intensity remains unchanged by enforcing a zero mean on P : $\hat{P} = P - \frac{1}{N} \sum_p P(p)$

- To remove undesirable boundaries from \hat{P} we apply a bilateral filter
- We create a map $M = e^{(\beta \hat{P})}$
- Our modifier is defined by: $I_{\text{regional}} = \frac{1}{N} \sum_i M I_i$

We observe that if:

- $\beta = 0$ the result remains unchanged
- $\beta > 0$ emphasizes regions where $\hat{P} > 0$
- $\beta < 0$ emphasizes regions where $\hat{P} < 0$

Implementation notes To get the intensity image \bar{I}_i we convert our BGR image to HSV color space. Then we extract the V component and work with that as our intensity image. Once we have done our calculations we return the V component to our HSV image, and convert it back to a BGR image.

Results In figure 10 we applied the regional light modifier to image 003.png from the cafe image set. The right hand side of this image is well lit, while the left hand side is almost completely dark. The results show that for $\beta = 0$ we get the original image back, for $\beta > 0$, the bright regions are emphasized, and for $\beta < 0$ the bright regions are de-emphasized.



Figure 10: Regional light modifier applied to the cafe fill light image.

Top Left: $\beta = 0$, **TopRight :** $\beta > 0$
Bottom : $\beta < 0$

6 Future work

An important way in which we could improve the speed of our implementation is to make use of GPU acceleration. Ideally we would like to have a library that allows NumPy to do calculations directly on the GPU. A brief search revealed a library called [Theano](#)¹⁵ that appears to suit our needs.

Another feature that would make our implementation more approachable to the general public would be to add a graphical user interface. Since this is a proof of concept we deemed this outside the scope of the project.

¹⁵<http://wwwdeeplearningnetsoftwaretheano/>